



# Inside The Indexer

How Clair V4 extracts and persists the contents of containers.

By: Louis DeLosSantos  
Principal Software Engineer, Clair

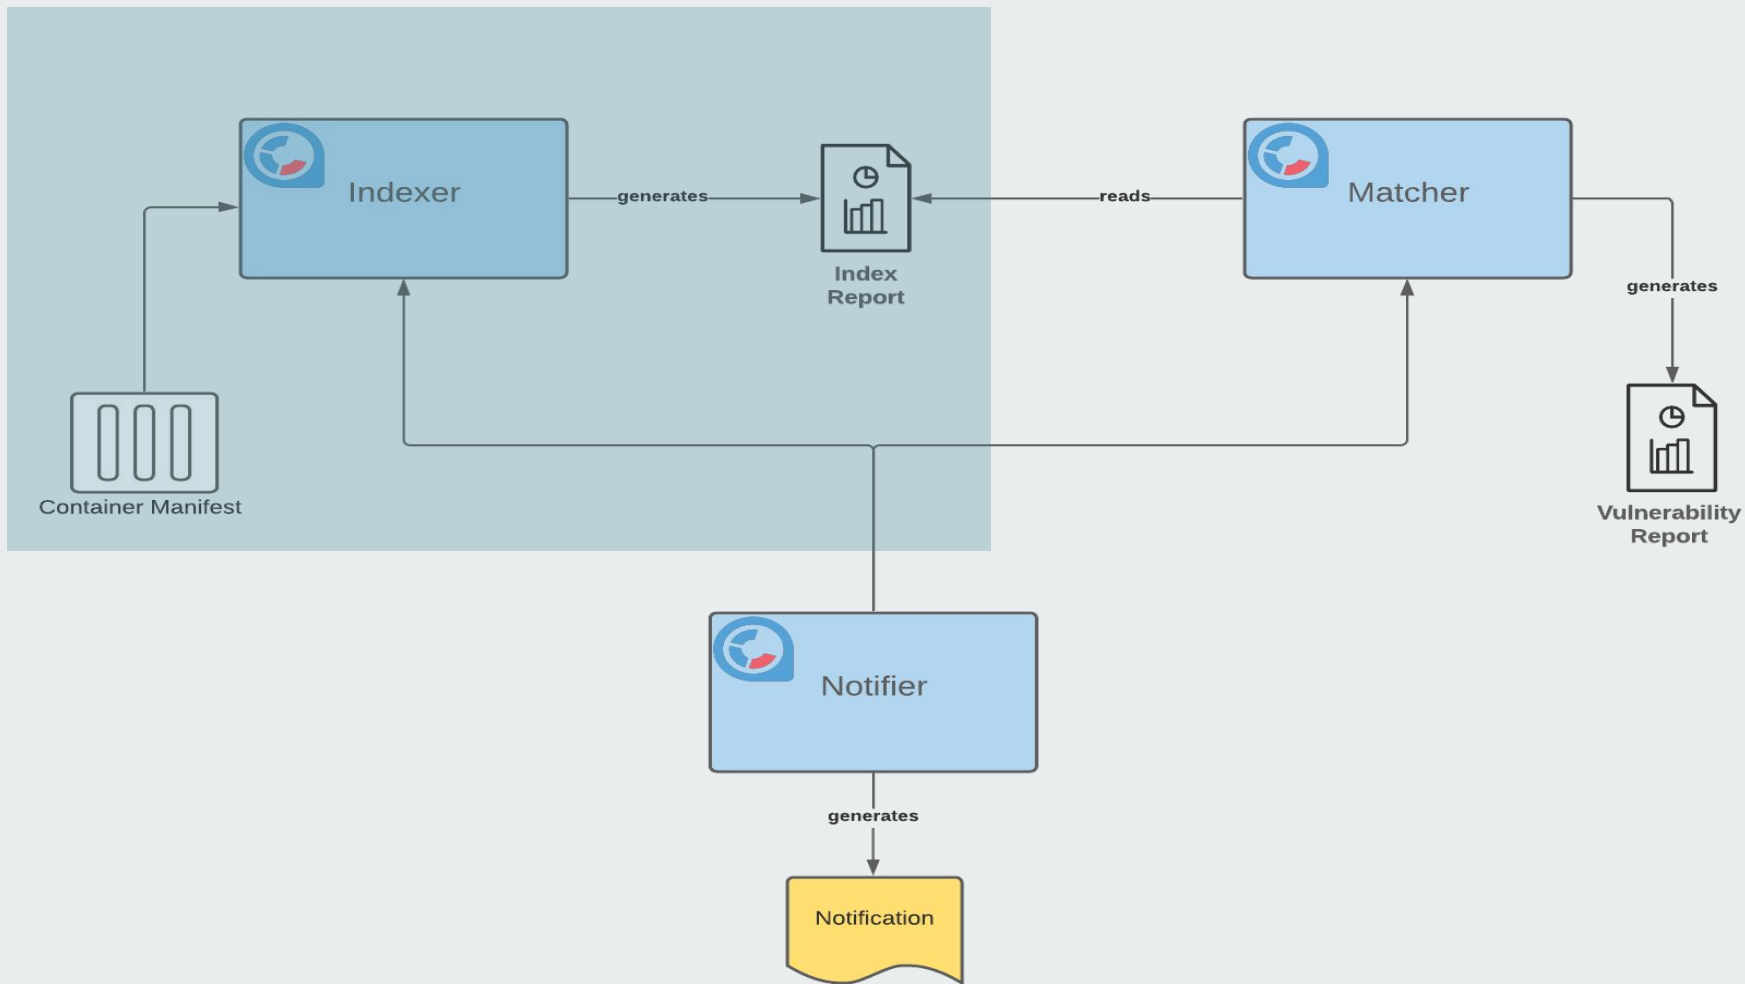


# Introduction



# What is "Indexing"

- First step in Clair's analysis pipeline
- Responsible for Index Report creation





# Key Components

# Manifests

```
package claircore

// Manifest represents a container image. Layers array MUST be indexed
// in the order that image layers are stacked.
type Manifest struct {
    // content addressable hash. should be able to be computed via
    // the hashes of all included layers
    Hash Digest `json:"hash"`
    // an array of filesystem layers indexed in the same order as the cooresponding image
    Layers []*Layer `json:"layers"`
}
```

# Index Reports

```
// IndexReport provides a database for discovered artifacts in an image.
//
// IndexReports make heavy usage of lookup maps to associate information
// without repetition.
type IndexReport struct {
    // the manifest hash this IndexReport is describing
    Hash Digest `json:"manifest_hash"`
    // the current state of the index operation
    State string `json:"state"`
    // all discovered packages in this manifest key'd by package id
    Packages map[string]*Package `json:"packages"`
    // all discovered distributions in this manifest key'd by distribution id
    Distributions map[string]*Distribution `json:"distributions"`
    // all discovered repositories in this manifest key'd by repository id
    Repositories map[string]*Repository `json:"repository"`
    // a list of environment details a package was discovered in key'd by package id
    Environments map[string][]*Environment `json:"environments"`
    // whether the index operation finished successfully
    Success bool `json:"success"`
    // an error string in the case the index did not succeed
    Err string `json:"err"`
}
```

# Scanner Interfaces

```
// PackageScanner provides an interface for unique identification or a PackageScanner
// and a Scan method for extracting installed packages from an individual container layer
type PackageScanner interface {
    VersionedScanner
    // Scan performs a package scan on the given layer and returns all
    // the found packages
    Scan(context.Context, *claircore.Layer) ([]*claircore.Package, error)
}

type DistributionScanner interface {
    VersionedScanner
    Scan(context.Context, *claircore.Layer) ([]*claircore.Distribution, error)
}

type RepositoryScanner interface {
    VersionedScanner
    Scan(context.Context, *claircore.Layer) ([]*claircore.Repository, error)
}
```



# Coalescer

```
// layerArtifact aggregates the any artifacts found within a layer
type LayerArtifacts struct {
    Hash    claircore.Digest
    Pkgs    []*claircore.Package
    Dist    []*claircore.Distribution // each layer can only have a single distribution
    Repos   []*claircore.Repository
}

// Coalescer takes a set of layers and creates coalesced IndexReport.
//
// A coalesced IndexReport should provide only the packages present in the
// final container image once all layers were applied.
type Coalescer interface {
    Coalesce(ctx context.Context, artifacts []*LayerArtifacts) (*claircore.IndexReport, error)
}
```

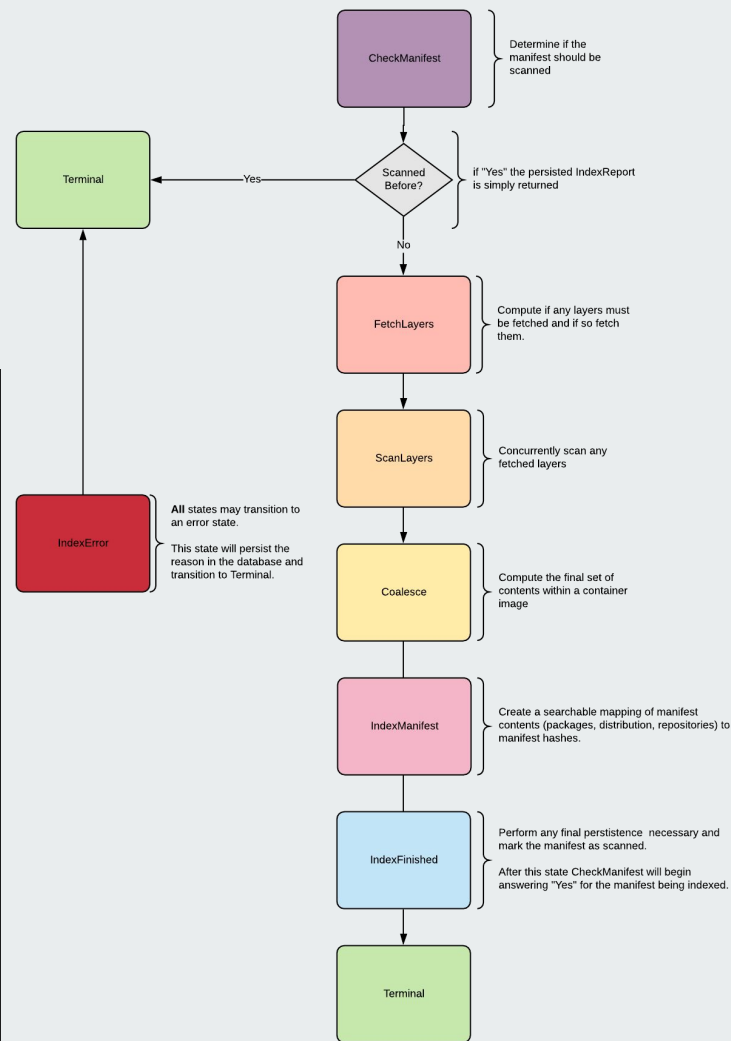


# Architecture

- RESTful HTTP API
- Finite State Machine

# Finite State Machine

```
// run executes each stateFunc and blocks until either an error occurs or
// a Terminal state is encountered.
func (s *Controller) run(ctx context.Context) {
    state, err := stateToStateFunc[s.getState()](ctx, s)
    if err != nil {
        s.handleError(ctx, err)
        return
    }
    if state == Terminal {
        return
    }
    s.setState(state)
    err = s.Store.SetIndexReport(ctx, s.report)
    if err != nil {
        s.handleError(ctx, err)
        return
    }
    s.run(ctx)
}
```



CheckManifest

Determine if the  
manifest should be  
scanned

604ac6bb

registry.redhat.io/ubi8/ubi-minimal

e84bb6b9

df08081

cc40570f



604ac6bb ✓

registry.redhat.io/ubi8/ubi-minimal

FetchLayers

Compute if any layers must  
be fetched and if so fetch  
them.

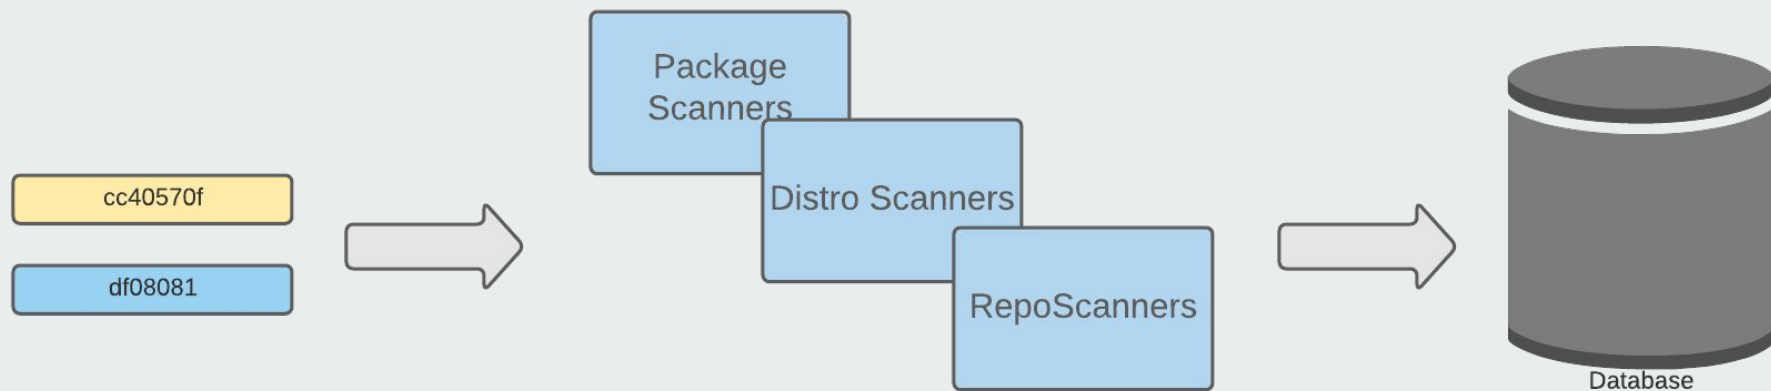



blob fetch



Buffer layers to disk





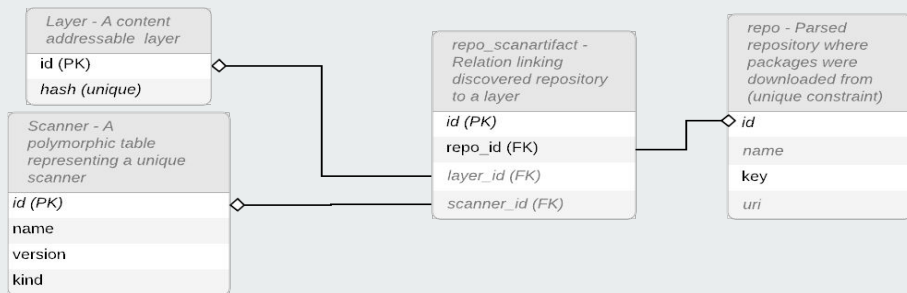
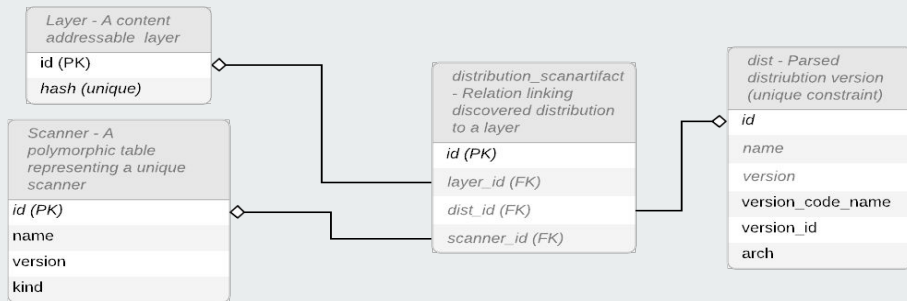
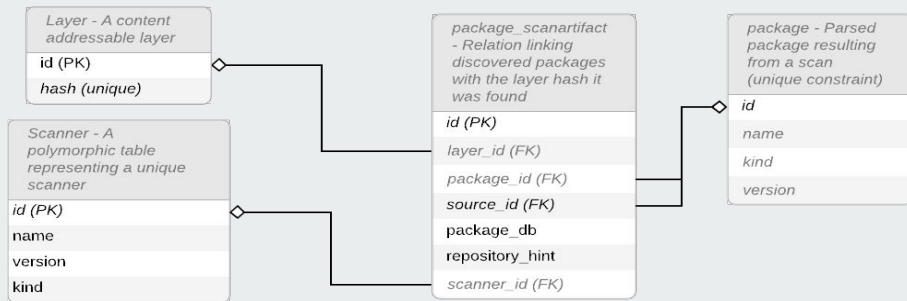
# Scanner Interfaces

```
// PackageScanner provides an interface for unique identification or a PackageScanner
// and a Scan method for extracting installed packages from an individual container layer
type PackageScanner interface {
    VersionedScanner
    // Scan performs a package scan on the given layer and returns all
    // the found packages
    Scan(context.Context, *claircore.Layer) ([]*claircore.Package, error)
}

type DistributionScanner interface {
    VersionedScanner
    Scan(context.Context, *claircore.Layer) ([]*claircore.Distribution, error)
}

type RepositoryScanner interface {
    VersionedScanner
    Scan(context.Context, *claircore.Layer) ([]*claircore.Repository, error)
}
```

# Scan Artifacts



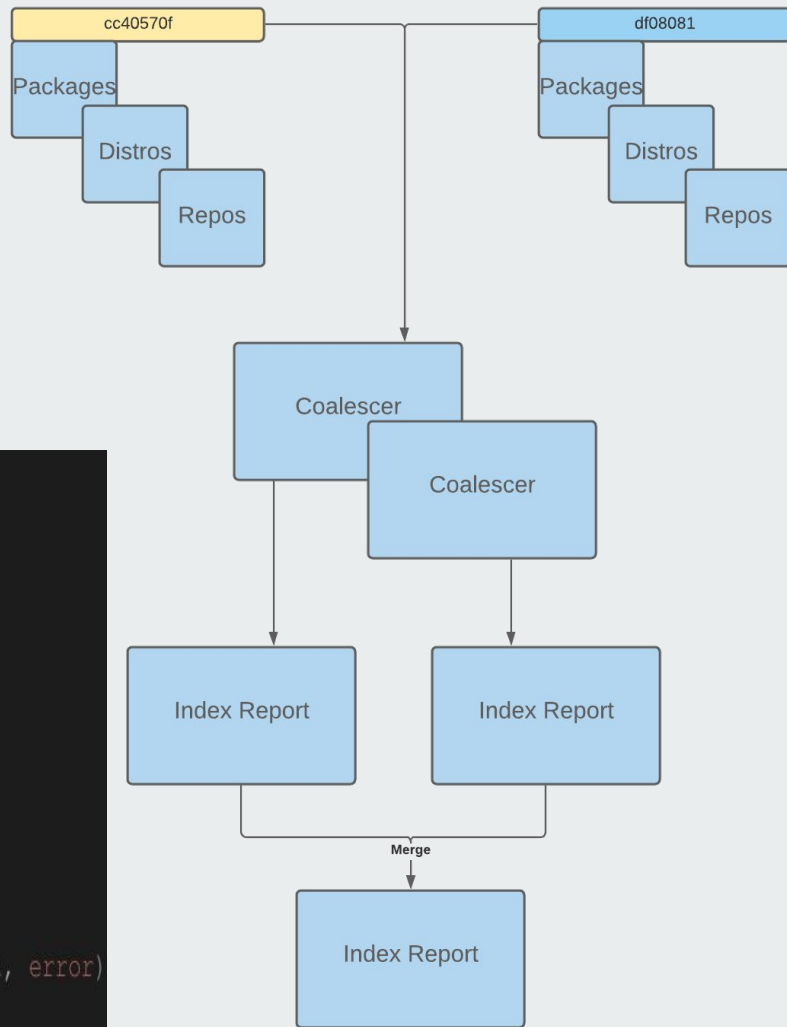


Coalesce

Compute the final set of  
contents within a container  
image

```
// layerArtifact aggregates the any artifacts found within a layer
type LayerArtifact struct {
    Hash    claircore.Digest
    Pkgs    []*claircore.Package
    Dist    []*claircore.Distribution // each layer can only have a single distribution
    Repos   []*claircore.Repository
}

// Coalescer takes a set of layers and creates coalesced IndexReport.
//
// A coalesced IndexReport should provide only the packages present in the
// final container image once all layers were applied.
type Coalescer interface {
    Coalesce(ctx context.Context, artifacts []*LayerArtifact) (*claircore.IndexReport, error)
}
```



IndexManifest

Create a searchable mapping of manifest contents (packages, distribution, repositories) to manifest hashes.

Manifest - A content addressable manifest	
id (PK)	
hash (unique)	

manifest_index - A relation providing the artifacts discoverable for a coalesced manifest.	
id (PK)	
package_id (FK)	
dist_id (FK)	
repo_id (FK)	
manifest_id (FK)	

package - Parsed package resulting from a scan (unique constraint)	
id	
name	
kind	
version	

dist - Parsed distribution version (unique constraint)	
id	
name	
version	
version_code_name	
version_id	
arch	

repo - Parsed repository where packages were downloaded from (unique constraint)	
id	
name	
key	
uri	


IndexFinished

Perform any final persistence necessary and mark the manifest as scanned.

After this state CheckManifest will begin answering "Yes" for the manifest being indexed.



IndexReport  
604ac6bb

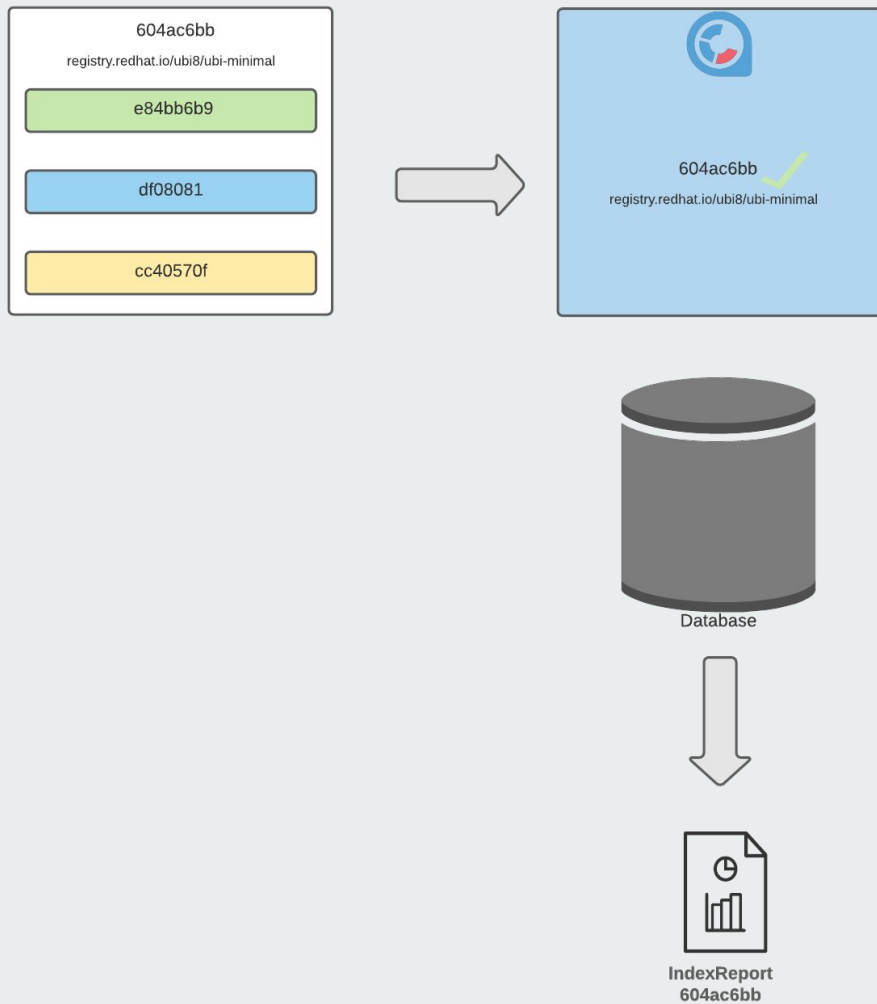


Database

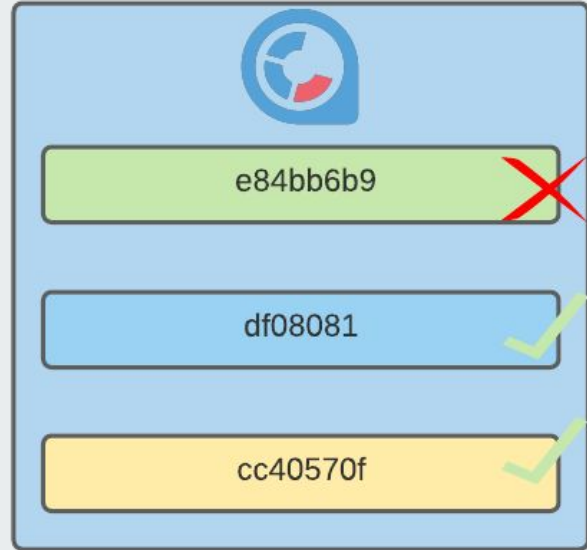


# Deferring Work

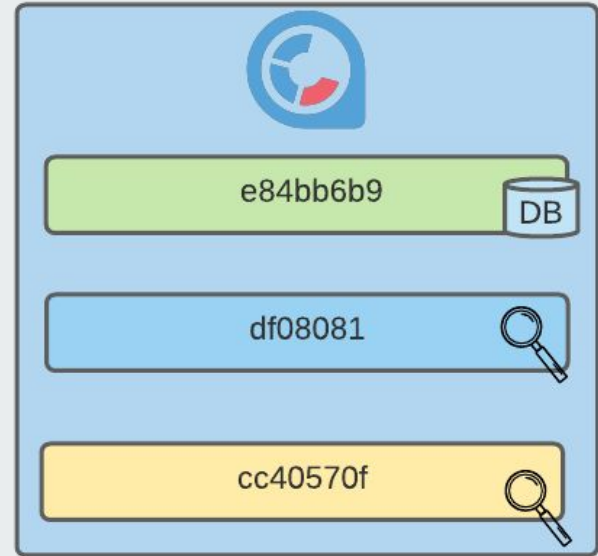
# Manifest Seen



# Determine layers to scan



# Scan only necessary layers





# Info

- [ldelossa@redhat.com](mailto:ldelossa@redhat.com)
- <https://github.com/quay/clair>
- <https://github.com/quay/claircore/>